Haoyi Zhu

✉ hyizhu1108@gmail.com | 🔗 https://haoyizhu.site/ | ⌂ Haoyi Zhu | 🎓 Google Scholar

EDUCATION

University of Science and Technology of China (USTC), China   *Sep. 2023 – Present*
Ph.D. in Computer Science (Advisor: Prof. Xiaogang Wang)
Shanghai Jiao Tong University (SJTU), China   *Sep. 2019 – Jun. 2023*
B.E. in Artificial Intelligence Honor Class (Advisor: Prof. Cewu Lu)

SELECTED AWARDS

Outstanding Paper Award, ICCV 2025 RIWM Workshop   *Oct. 2025*
National Scholarship   *Oct. 2025*
First Price Academic Scholarship, USTC   *Sep. 2025*
XingQi Intern, Shanghai AI Lab **(6/160+)**   *Apr. 2025*
Outstanding Paper Award, NeurIPS 2022   *Dec. 2022*

INTERNSHIPS

Shanghai AI Lab   *Sep. 2023 - Nov. 2025*
Research Intern (Advisor: Prof. Tong He & Prof. Wanli Ouyang)
- Foundation World Model
- General 3D Embodied Representation Learning

MVIG, Shanghai Jiao Tong University (SJTU)   *Sep. 2019 - Jun. 2023*
Research Intern (Advisor: Dr. Hao-Shu Fang & Prof. Cewu Lu)
- Human Pose Estimation
- General Robot Manipulation

GEAR Lab, NVIDIA   *Feb. 2022 - Feb. 2023*
Research Intern (Advisor: Dr. Guanzhi Wang & Dr. Jim Fan)
- General Open-Ended Embodied Agent

Xu Lab, Carnegie Mellon University (CMU)   *Apr. 2021 – Feb. 2022*
Research Intern (Advisor: Prof. Min Xu)
- AI for Biology

PUBLICATIONS   **Total Citations: 2065**

Journal Papers

PonderV2: Pave the Way for 3D Foundation Model with A Universal Pre-training Paradigm [paper] [code]
Haoyi Zhu*, Honghui Yang*, Xiaoyang Wu*, Di Huang*, Sha Zhang, Xianglong He, Tong He, Hengshuang Zhao, Chunhua Shen, Yu Qiao, Wanli Ouyang
*IEEE Transactions on Pattern Analysis and Machine Intelligence (**TPAMI**), 2025*

AlphaTracker: A Multi-Animal Tracking and Behavioral Analysis Tool [paper] [code]
Zexin Chen, Ruihan Zhang, Hao-Shu Fang, Yu E Zhang, Aneesh Bal, Haowen Zhou, Rachel R Rock, Nancy Padilla-Coreano, Laurel R Keyes, **Haoyi Zhu**, Yong-Lu Li, Takaki Komiyama, Kay M Tye, Cewu Lu
Frontiers in Behavioral Neuroscience, 2023

AlphaPose: Whole-Body Regional Multi-Person Pose Estimation and Tracking in Real-Time [paper] [code]
Hao-Shu Fang, Jiefeng Li, Hongyang Tang, Chao Xu*, **Haoyi Zhu***, Yuliang Xiu, Yong-Lu Li, Cewu Lu
*[>8.3K GitHub Stars] IEEE Transactions on Pattern Analysis and Machine Intelligence (**TPAMI**), 2022*

Conference Papers (Selected)

OmniWorld: A Multi-Domain and Multi-Modal Dataset for 4D World Modeling [paper] [website] [code]
..., Haoyi Zhu, ... (19 authors)
arXiv preprint, 2025

WinT3R: Window-Based Streaming Reconstruction with Camera Token Pool [paper] [website] [code]
Yifan Wang, Jianjun Zhou, **Haoyi Zhu**, Wenzheng Chang, Yang Zhou, Zizun Li, Junyi Chen, Jiangmiao Pang, Chunhua Shen, Tong He
arXiv preprint, 2025

π^3**: Scalable Permutation-Equivariant Visual Geometry Learning** [paper] [website] [code]
Yifan Wang, Jianjun Zhou, **Haoyi Zhu**, Wenzheng Chang, Yang Zhou, Zizun Li, Junyi Chen, Jiangmiao Pang, Chunhua Shen, Tong He
arXiv preprint, 2025

DeepVerse: 4D Autoregressive Video Generation as a World Model [paper] [website] [code]
Junyi Chen, **Haoyi Zhu**, Xianglong He, Yifan Wang, Jianjun Zhou, Wenzheng Chang, Yang Zhou, Zizun Li, Zhoujie Fu, Jiangmiao Pang, Tong He
arXiv preprint, 2025

VQ-VLA: Improving Vision-Language-Action Models via Scaling Vector-Quantized Action Tokenizers [paper] [website][code]
Yating Wang, **Haoyi Zhu**, Mingyu Liu, Jiange Yang, Hao-Shu Fang, Tong He
*International Conference on Computer Vision (**ICCV**), 2025*

Aether: Geometric-Aware Unified World Modeling [paper] [website] [code]
Haoyi Zhu, Yifan Wang, Jianjun Zhou, Wenzheng Chang, Yang Zhou, Zizun Li, Junyi Chen, Chunhua Shen, Jiangmiao Pang, Tong He
*International Conference on Computer Vision (**ICCV**), 2025*
*[Outstanding Paper Award & Oral Presentation] ICCV 2025 Workshop on Reliable and Interactive World Models (**RIWM**)*

CoMo: Learning Continuous Latent Motion from Internet Videos for Scalable Robot Learning [paper]
Jiange Yang, Yansong Shi, **Haoyi Zhu**, Mingyu Liu, Kaijing Ma, Yating Wang, Gangshan Wu, Tong He, Liming Wang
arXiv preprint, 2025

SPA: 3D Spatial-Awareness Enables Effective Embodied Representation [paper] [website] [code]
Haoyi Zhu, Honghui Yang, Yating Wang, Jiange Yang, Limin Wang, Tong He
*International Conference on Learning Representations (**ICLR**), 2025*

Tra-MoE: Learning Trajectory Prediction Model from Multiple Domains for Adaptive Policy Conditioning [paper] [code]
Jiange Yang, **Haoyi Zhu**, Yating Wang, Gangshan Wu, Tong He, Limin Wang
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR**), 2025*

Point Cloud Matters: Rethinking the Impact of Different Observation Spaces on Robot Learning [paper] [website] [code]
Haoyi Zhu, Yating Wang, Di Huang, Weicai Ye, Wanli Ouyang, Tong He
*Advances in Neural Information Processing Systems (**NeurIPS**), 2024*

RH20T: A Comprehensive Robotic Dataset for Learning Diverse Skills in One-Shot [paper] [website] [code]
Hao-Shu Fang, Hongjie Fang, Zhenyu Tang, Jirong Liu, Chenxi Wang, Junbo Wang, **Haoyi Zhu**, Cewu Lu
*IEEE International Conference on Robotics and Automation (**ICRA**), 2024*

UniPAD: A Universal Pre-Training Paradigm for Autonomous Driving [paper] [code]
Honghui Yang, Sha Zhang, Di Huang, Xiaoyang Wu, **Haoyi Zhu**, Tong He, Shixiang Tang, Hengshuang Zhao, Qibo Qiu, Binbin Lin, Xiaofei He, Wanli Ouyang
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR**), 2024*

X-NeRF: Explicit Neural Radiance Field for Multi-Scene 360° Insufficient RGB-D Views [paper] [code]
Haoyi Zhu, Hao-Shu Fang, Cewu Lu
*IEEE/CVF Winter Conference on Applications of Computer Vision (**WACV**), 2023*

MineDojo: Building Open-Ended Embodied Agents with Internet-Scale Knowledge [paper] [website] [code]
Linxi Fan, Guanzhi Wang, Yunfan Jiang, Ajay Mandlekar, Yuncong Yang, **Haoyi Zhu**, Andrew Tang, De-An Huang, Yuke Zhu, Anima Anandkumar
*[Outstanding Paper Award] Advances in Neural Information Processing Systems (**NeurIPS**), 2022*

For the full publication list, please visit my Google Scholar page.

SELECTED PROJECTS

Aether
Scalable and Generalizable Foundation World Model *Nov. 2024 – Present*

- Proposed **Aether**, a geometry-aware unified world model. **Aether** features 3 core capabilities: (1) 4D dynamic reconstruction, (2) action-conditioned video prediction, and (3) goal-conditioned visual planning. Building upon video generation models, our framework demonstrates unprecedented synthetic-to-real generalization

despite *never* observing real-world data during training. Won the **Outstanding Paper Award** at ICCV 2025 RIWM workshop.
- Developed **DeepVerse**, a 4D auto-regressive world model. **DeepVerse** can fantasize the entire world behind images and enable free exploration through interaction. It demonstrates exceptional capabilities in modeling dynamics, geometry, and physical laws.
- Introduced π^3, a novel feed-forward neural network that revolutionizes visual geometry reconstruction by eliminating the need for a fixed reference view. With a fully permutation-equivariant architecture, π^3 achieves SOTA performance on a wide range of tasks, including camera pose estimation, monocular/video depth estimation, and dense point map estimation.
- Mentored several undergraduate and junior Ph.D. students on robot learning topics, including action tokenizers, cross-embodiment policies, point cloud policies, *etc.*

SPA
Spatial intelligence and General 3D Embodied Representation Learning   *Mar. 2023 – Oct. 2024*

- Conducted extensive research on general 3D representation for embodied AI. Published three first-author papers.
- Proposed a universal 3D pre-training framework, ***PonderV2***, which leverages differentiable neural rendering to learn point cloud representations. Achieved state-of-the-art (SOTA) performance across 11 indoor and outdoor benchmarks.
- Performed an in-depth investigation of observation spaces in robot learning, culminating in the discovery of the critical importance of ***Point Cloud Matters***.
- Developed ***SPA***, an innovative representation learning framework that highlights the significance of 3D spatial awareness in embodied AI. Delivered the most comprehensive evaluation of embodied representation learning to date, surpassing over 10 SOTA methods.

MineDojo
Building Open-Ended Embodied Agents with Internet-Scale Knowledge   *Feb. 2022 – Jul. 2022*

- Contributed to the ***MineDojo*** project, a new framework designed to build generally capable agents with internet-scale knowledge in Minecraft. Won the **Outstanding Paper Award** at NeurIPS 2022.
- Co-developed the YouTube database, which comprises over 730,000 narrated Minecraft videos, totaling approximately 300,000 hours of content and 2.2 billion English transcripts.
- Constructed MineCLIP, a video CLIP foundation model that acts as a learned reward function, enabling agents to solve diverse open-ended tasks specified in natural language.
- Authored APIs for three databases (YouTube, Minecraft Wiki, and Reddit), along with the corresponding documentation and part of the official project website.

ACADEMIC SERVICES

Reviewer for journals, including IJCV, TMLR, *etc*.
Reviewer for conferences, including ICCV, CVPR, NeurIPS, ICLR, WACV, ICIP, *etc*.

INVITED TALKS

- 2025/08/26, Prof. Lin Shao's Lab, NUS, "World Model and Spatial Intelligence"
- 2025/06/12, 3D CV Workshop (3D视觉工坊), "Towards Generalizable and Scalable Spatial Embodied Intelligence" [link]
- 2025/05/22, Embodied Intelligent Mind (具身智能之心), "Towards Generalizable and Scalable Spatial Embodied Intelligence"
- 2025/04/22, Prof. Yanyong Zhang's Lab, USTC, "Towards Scalable World Model for General Embodied AI"
- 2025/03/27, Graduate Academic Forum, USTC, "Towards Spatial Intelligence: From 3D Vision to Embodied AI"
- 2024/12/27, ZhiXingXing Embodied AI Frontier Lecture (智猩猩具身智能前沿讲座), "Towards Spatial Intelligence: From 3D Vision to Embodied AI" [link]